Running a Job on a MapReduce v2 Cluster (2)

Slave Nodes

Client

Resource Manager(s)

Job History Server

NodeManager — Map Task 1 — DataNode — Block1

NodeManager — Map Task 2 / Reduce Task 1 — DataNode — Block2

NodeManager — **MapReduce Application Master** — DataNode

NodeManager — Reduce Task 2 — DataNode

HDFS: `mydata`

Name Node(s)